

# Executable and Linkable Format

# ELF

- file format used in Linux for:
  - object files (gcc -c file.c)
  - **executable** (gcc -o file file.c)
  - shared libraries (gcc -shared -o file.so file.c)
  - core dumps
- program header table (PHT) describe *segments*, which contains one or more *sections*
- *loadable segments* provide the *execution* view of the object file

```
→ stack git:(master) X file ./stack4
./stack4: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV), dynamically linked (uses shared
dID[sha1]=87beb67a2e7ad1edb7bf56520d4bbd2e07827c88, stripped
→ stack git:(master) X xxd ./stack4 | head -n 1
00000000: 7f45 4c46 0101 0100 0000 0000 0000 0000  .ELF.....
```

- an executable begin a process when it is loaded in memory and executed
- Linux: *Executable and Linkable Format* (ELF): provides two interfaces: executable and linkable which are either described into the ELF header
- the *linkable* interface is not required for execution, therefore it will not be examined
- the *executable* interface is described by *program header* which are stored in a *program header table*

# Tools

- objdump
- readelf
- file

# print PHT using *readelf*

```
→ stack git:(master) X readelf -l ./stack4
```

Elf file type is EXEC (Executable file)

Entry point 0x8048380

There are 9 program headers, starting at offset 52

to be loaded in memory

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x08048034	0x08048034	0x00120	0x00120	R E	0x4
INTERP	0x000154	0x08048154	0x08048154	0x00013	0x00013	R	0x1
[Requesting program interpreter: /lib/ld-linux.so.2]							
LOAD	0x000000	0x08048000	0x08048000	0x00660	0x00660	R E	0x1000
LOAD	0x000f08	0x08049f08	0x08049f08	0x00120	0x00124	RW	0x1000
DYNAMIC	0x000f14	0x08049f14	0x08049f14	0x000e8	0x000e8	RW	0x4
NOTE	0x000168	0x08048168	0x08048168	0x00044	0x00044	R	0x4
GNU EH FRAME	0x000584	0x08048584	0x08048584	0x0002c	0x0002c	R	0x4
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RWE	0x10
GNU_RELRO	0x000f08	0x08049f08	0x08049f08	0x000f8	0x000f8	R	0x1

Section to Segment mapping:

Segment Sections...

```
00
01      .interp
02      .interp .note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rel.dyn .rel.plt .init .plt .text .fini .rodata .eh_frame_hdr .eh_frame
03      .init_array .fini_array .jcr .dynamic .got .got.plt .data .bss
04      .dynamic
05      .note.ABI-tag .note.gnu.build-id
06      .eh_frame_hdr
07
08      .init_array .fini_array .jcr .dynamic .got
```

# PHT `aka` segment header table

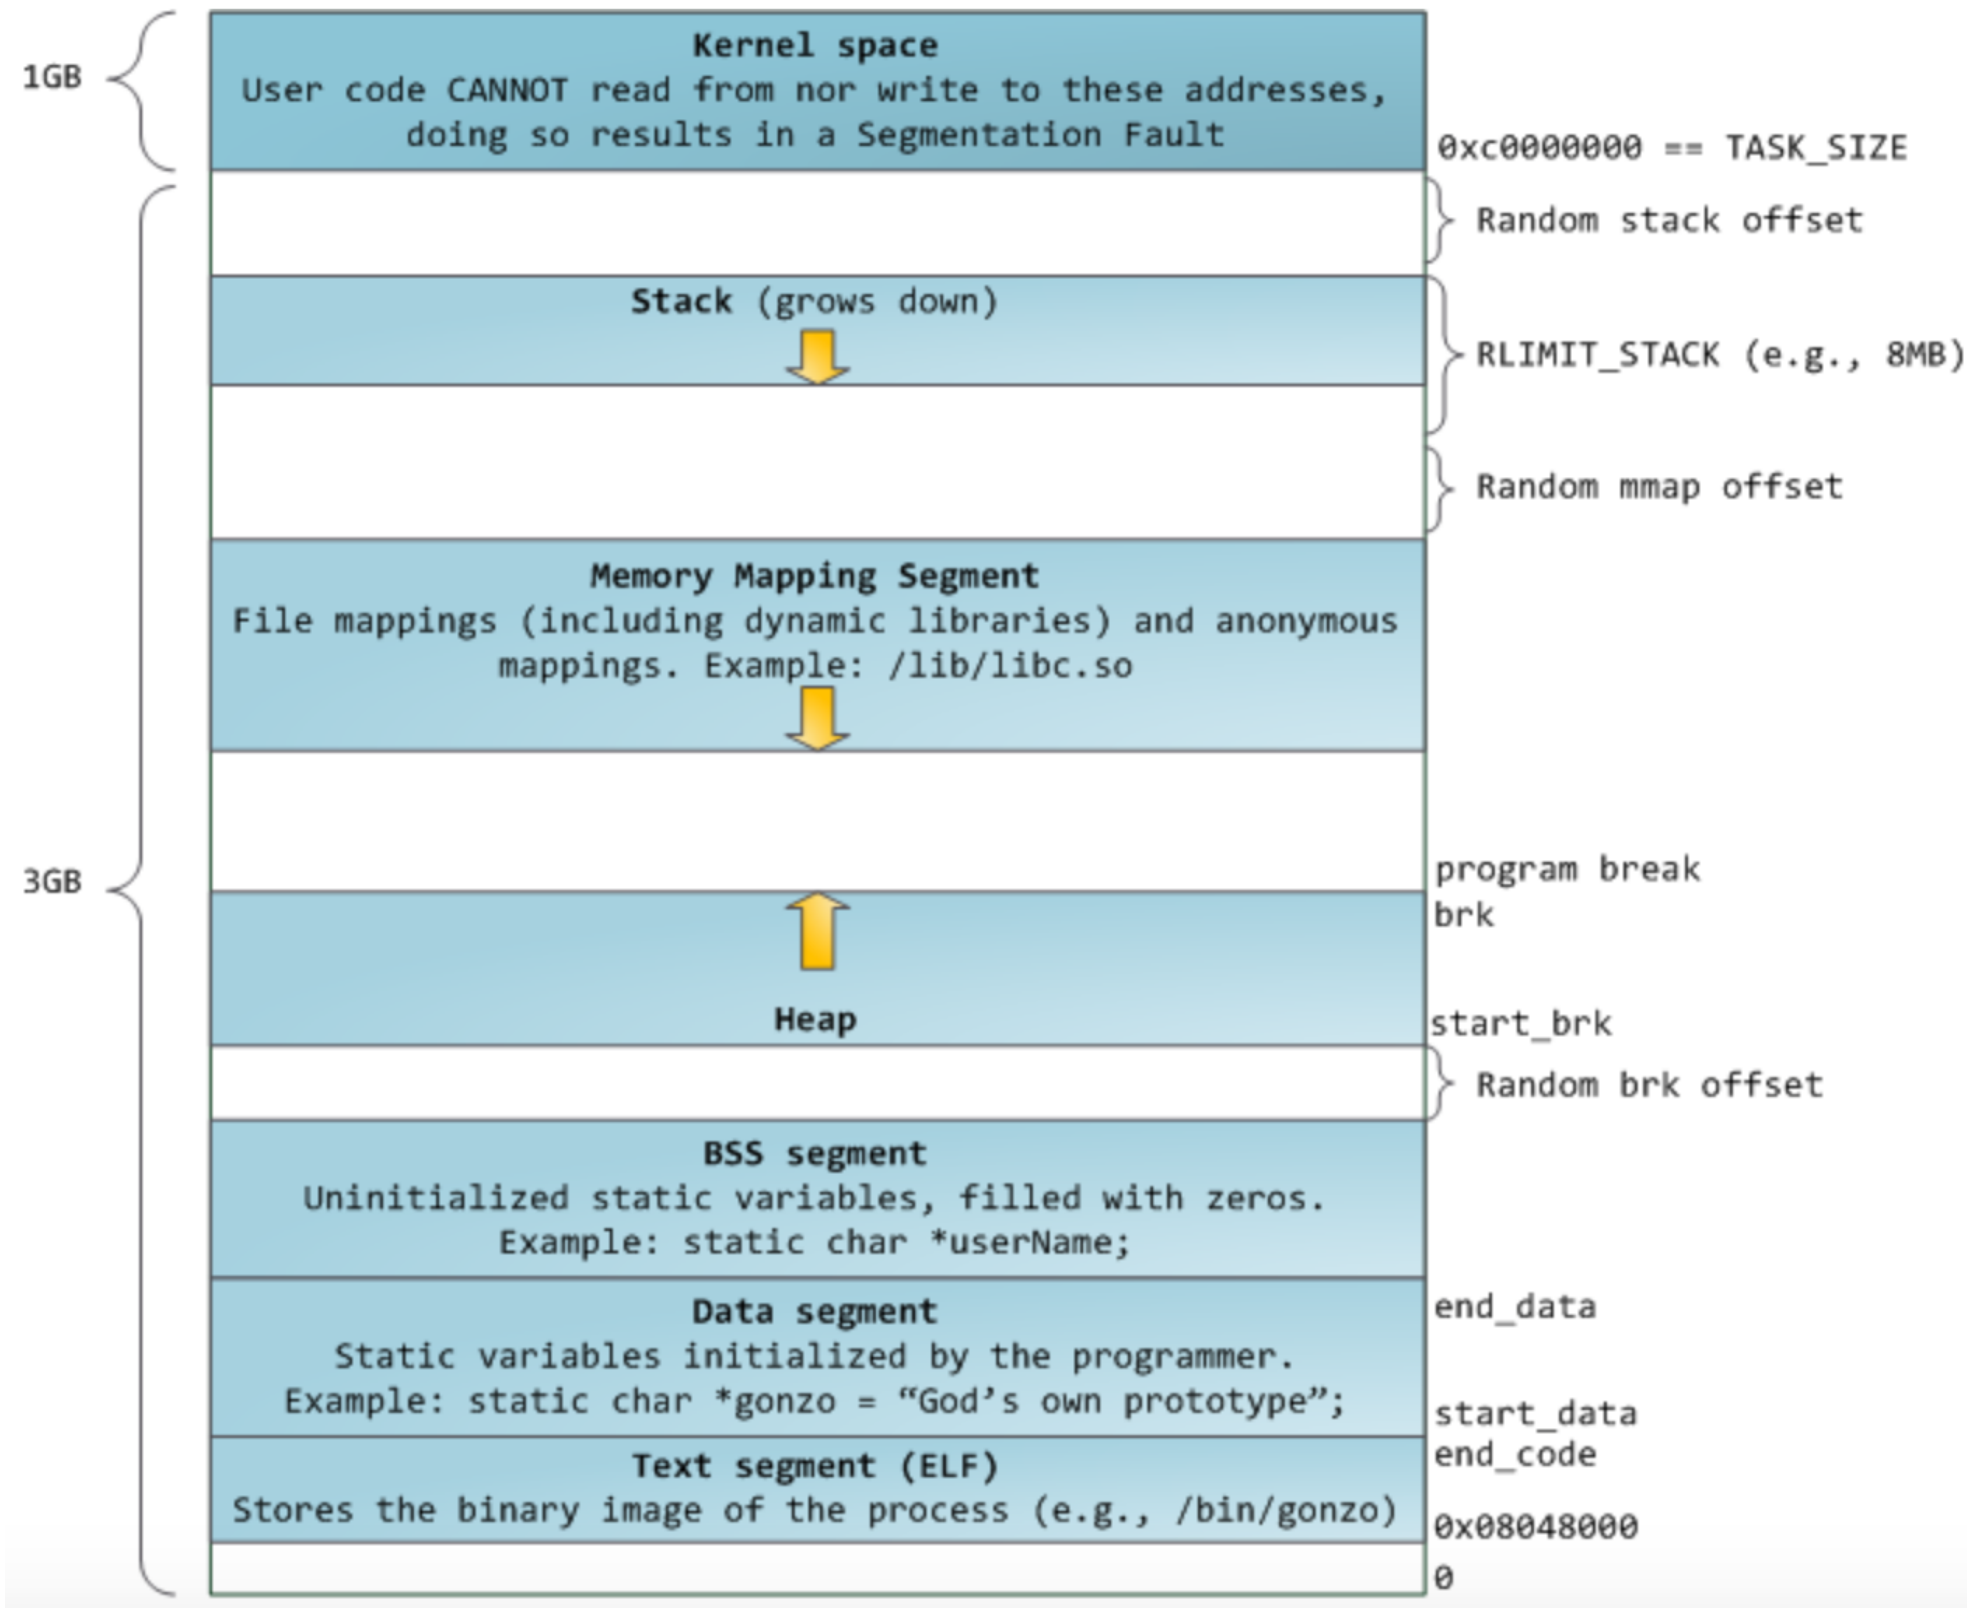
- to load a file in memory, program headers are used to provide informations like:
  - parts of file to be loaded into memory at runtime
  - locations of important data at runtime
- PHT describe *segments* as an offset from the start of the file and a size
- executable and shared object contains *segments*
- A *segments* is grouping of one or more sections

- neither SHT nor PHT have fixed address in memory
- to locate them we use the ELF header

```
→ stack git:(master) X readelf -h ./stack4
ELF Header:
  Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF32
  Data:                                   2's complement, little endian
  Version:                             1 (current)
  OS/ABI:                               UNIX - System V
  ABI Version:                         0
  Type:                                 EXEC (Executable file)
  Machine:                              Intel 80386
  Version:                              0x1
  Entry point address:                  0x8048380
  Start of program headers:              52 (bytes into file)
  Start of section headers:              4460 (bytes into file)
  Flags:                                0x0
  Size of this header:                   52 (bytes)
  Size of program headers:               32 (bytes)
  Number of program headers:              9
  Size of section headers:               40 (bytes)
  Number of section headers:              28
  Section header string table index: 27
```

- Process in memory are divided into three regions:
  - text: fixed by the program and includes code (instructions) and read-only data. loaded at fixed address, for i386 binary is 0x80480000 (except for Position Independent Executable)
  - data: contains initialized and uninitialized data (BSS) and static variables. it's size can be changed with the `brk(2)` system call
  - stack: we already know this data structure...





# VMA

- In Linux, a process' linear address space is organized in sets of virtual memory areas (VMAs)
- An object file's loadable segment corresponds to at least one VMAs
- The output from `/proc/<pid>/maps` lists the memory areas in this process's address space:

```
→ stack git:(master) X cat /proc/self/maps
00400000-0040b000 r-xp 00000000 08:01 524314      /bin/cat
0060a000-0060b000 r--p 0000a000 08:01 524314      /bin/cat
0060b000-0060c000 rw-p 0000b000 08:01 524314      /bin/cat
0113b000-0115c000 rw-p 00000000 00:00 0          [heap]
7fce195ac000-7fce19c8e000 r--p 00000000 08:01 660521      /usr/lib/locale/locale-archive
7fce19c8e000-7fce19e49000 r-xp 00000000 08:01 399528      /lib/x86_64-linux-gnu/libc-2.19.so
7fce19e49000-7fce1a048000 ---p 001bb000 08:01 399528      /lib/x86_64-linux-gnu/libc-2.19.so
7fce1a048000-7fce1a04c000 r--p 001ba000 08:01 399528      /lib/x86_64-linux-gnu/libc-2.19.so
7fce1a04c000-7fce1a04e000 rw-p 001be000 08:01 399528      /lib/x86_64-linux-gnu/libc-2.19.so
7fce1a04e000-7fce1a053000 rw-p 00000000 00:00 0
7fce1a053000-7fce1a076000 r-xp 00000000 08:01 399504      /lib/x86_64-linux-gnu/ld-2.19.so
7fce1a259000-7fce1a25c000 rw-p 00000000 00:00 0
7fce1a273000-7fce1a275000 rw-p 00000000 00:00 0
7fce1a275000-7fce1a276000 r--p 00022000 08:01 399504      /lib/x86_64-linux-gnu/ld-2.19.so
7fce1a276000-7fce1a277000 rw-p 00023000 08:01 399504      /lib/x86_64-linux-gnu/ld-2.19.so
7fce1a277000-7fce1a278000 rw-p 00000000 00:00 0
7ffdbf50a000-7ffdbf52b000 rw-p 00000000 00:00 0          [stack]
7ffdbf530000-7ffdbf532000 r--p 00000000 00:00 0          [vvar]
7ffdbf532000-7ffdbf534000 r-xp 00000000 00:00 0          [vdso]
ffffffff600000-ffffffff601000 r-xp 00000000 00:00 0          [vsyscall]
```

# PLT and GOT

- come vengono richiamate le funzioni contenute nelle *shared library*?

```
→ stack git:(master) X ldd ./got
linux-gate.so.1 => (0xf76e8000)
libc.so.6 => /lib/i386-linux-gnu/libc.so.6 (0xf751e000)
/lib/ld-linux.so.2 (0xf76ea000)
```

- ldd mostra le *shared library* richieste dal binario *got*
- *Procedure Linkage Table (PLT)*
- *Global Offset Table (GOT)*

# PLT and GOT

```
→ stack git:(master) X cat got.c
main()
{
    printf("GOT! \n");
}
```


prima dell'esecuzione

chiamata a puts

```
gdb-peda$ disass main
Dump of assembler code for function main:
   0x0804841d <+0>:    push    ebp
   0x0804841e <+1>:    mov     ebp,esp
   0x08048420 <+3>:    sub     esp,0x4
   0x08048423 <+6>:    mov     DWORD PTR [esp],0x80484d0
   0x0804842a <+13>:   call    0x80482f0 <puts@plt>
   0x0804842f <+18>:    leave
   0x08048430 <+19>:    ret
End of assembler dump.
```

codice della funzione puts?

```
gdb-peda$ disass puts
Dump of assembler code for function puts@plt:
   0x080482f0 <+0>:    jmp     DWORD PTR ds:0x804a00c
   0x080482f6 <+6>:    push    0x0
   0x080482fb <+11>:   jmp     0x80482e0
End of assembler dump.
gdb-peda$ printf "0x%x\n", *0x804a00c
0x80482f6
```



# PLT and GOT

```
gdb-peda$ disass main
Dump of assembler code for function main:
0x0804841d <+0>:      push    ebp
0x0804841e <+1>:      mov     ebp,esp
0x08048420 <+3>:      sub     esp,0x4
0x08048423 <+6>:      mov     DWORD PTR [esp],0x80484d0
0x0804842a <+13>:     call    0x80482f0 <puts@plt>
0x0804842f <+18>:     leave
0x08048430 <+19>:     ret
```

Impostiamo un  
breakpoint  
sulla prima istruzione  
del main ed eseguiamo

ora analizziamo  
il contenuto della  
GOT.  
ora contiene  
l'indirizzo:  
0xf7e76650

```
gdb-peda$ x/i 0x080482f0
0x80482f0 <puts@plt>:      jmp     DWORD PTR ds:0x804a00c
gdb-peda$
0x80482f6 <puts@plt+6>:    push    0x0
gdb-peda$
0x80482fb <puts@plt+11>:   jmp     0x80482e0
gdb-peda$ printf "0x%x\n", *0x804a00c
0xf7e76650
gdb-peda$ disass 0xf7e76650
Dump of assembler code for function puts:
0xf7e76650 <+0>:      push    ebp
0xf7e76651 <+1>:      push    edi
0xf7e76652 <+2>:      push    esi
0xf7e76653 <+3>:      push    ebx
```



```
➔ ~ cat /proc/31384/maps
```

```
08048000-08049000 r-xp 00000000 08:01 316714 /home/r0x/lezioni/sicII/esercizi/stack/got
08049000-0804a000 r-xp 00000000 08:01 316714 /home/r0x/lezioni/sicII/esercizi/stack/got
0804a000-0804b000 rwxp 00001000 08:01 316714 /home/r0x/lezioni/sicII/esercizi/stack/got
f7581000-f7582000 rwxp 00000000 00:00 0
f7582000-f772a000 r-xp 00000000 08:01 556235 /lib/i386-linux-gnu/libc-2.19.so
f772a000-f772c000 r-xp 001a8000 08:01 556235 /lib/i386-linux-gnu/libc-2.19.so
f772c000-f772d000 rwxp 001aa000 08:01 556235 /lib/i386-linux-gnu/libc-2.19.so
f772d000-f7730000 rwxp 00000000 00:00 0
f7746000-f7749000 rwxp 00000000 00:00 0
f7749000-f774b000 r--p 00000000 00:00 0 [vvar]
f774b000-f774d000 r-xp 00000000 00:00 0 [vdso]
f774d000-f776d000 r-xp 00000000 08:01 556237 /lib/i386-linux-gnu/ld-2.19.so
f776d000-f776e000 r-xp 0001f000 08:01 556237 /lib/i386-linux-gnu/ld-2.19.so
f776e000-f776f000 rwxp 00020000 08:01 556237 /lib/i386-linux-gnu/ld-2.19.so
ffe3d000-ffe5e000 rwxp 00000000 00:00 0 [stack]
```

```
gdb-peda$ x/i 0x080482f0
0x080482f0 <puts@plt>:      jmp     DWORD PTR ds:0x804a00c
gdb-peda$
0x080482f6 <puts@plt+6>:    push    0x0
gdb-peda$
0x080482fb <puts@plt+11>:   jmp     0x80482e0
gdb-peda$ printf "0x%x\n", *0x804a00c
0xf7e76650
gdb-peda$ disass 0xf7e76650
Dump of assembler code for function puts:
0xf7e76650 <+0>:      push    ebp
0xf7e76651 <+1>:      push    edi
0xf7e76652 <+2>:      push    esi
0xf7e76653 <+3>:      push    ebx
```